



# RE<sup>x</sup>: A DEVELOPMENT PLATFORM AND ONLINE LEARNING APPROACH FOR RUNTIME EMERGENT SOFTWARE SYSTEMS

**Barry Porter**, Matthew Grieves, Roberto Rodrigues Filho and David Leslie

School of Computing and Communications

Department of Mathematics and Statistics

Lancaster University, UK

[b.f.porter@lancaster.ac.uk](mailto:b.f.porter@lancaster.ac.uk)



# MOTIVATION

- Modern software remains **highly complex** to design, implement, maintain & configure, particularly for dynamic environments
  - this causes **high development costs** and **under-performing code**
- The state of the art in solving this is **self-adaptive systems**, which exhibit some awareness of *self* and of *environment*
  - However, these approaches relate only to the **configuration** element of systems development, and are also very **manual** in their definition, by:
    - Designing the base system as a non-adaptive one
    - Deciding which points of that system should be adaptable
    - Writing rules to determine how and when adaptation happens



# CONCEPT OVERVIEW

- We propose a paradigm of **continuous self-assembly**, in which the *initial construction* of software and its *later adaptation* are one continuous machine-driven process



# CONCEPT OVERVIEW

- We propose a paradigm of **continuous self-assembly**, in which the *initial construction* of software and its *later adaptation* are one continuous machine-driven process

**Goal** (description)  
(unit tests)

?



*Start with a goal and a pool of behavior fragments*




# CONCEPT OVERVIEW

- We propose a paradigm of **continuous self-assembly**, in which the *initial construction* of software and its *later adaptation* are one continuous machine-driven process

Goal (description)  
(unit tests)

?

Goal (description)  
(unit tests)



Start with a goal and a pool of behavior fragments

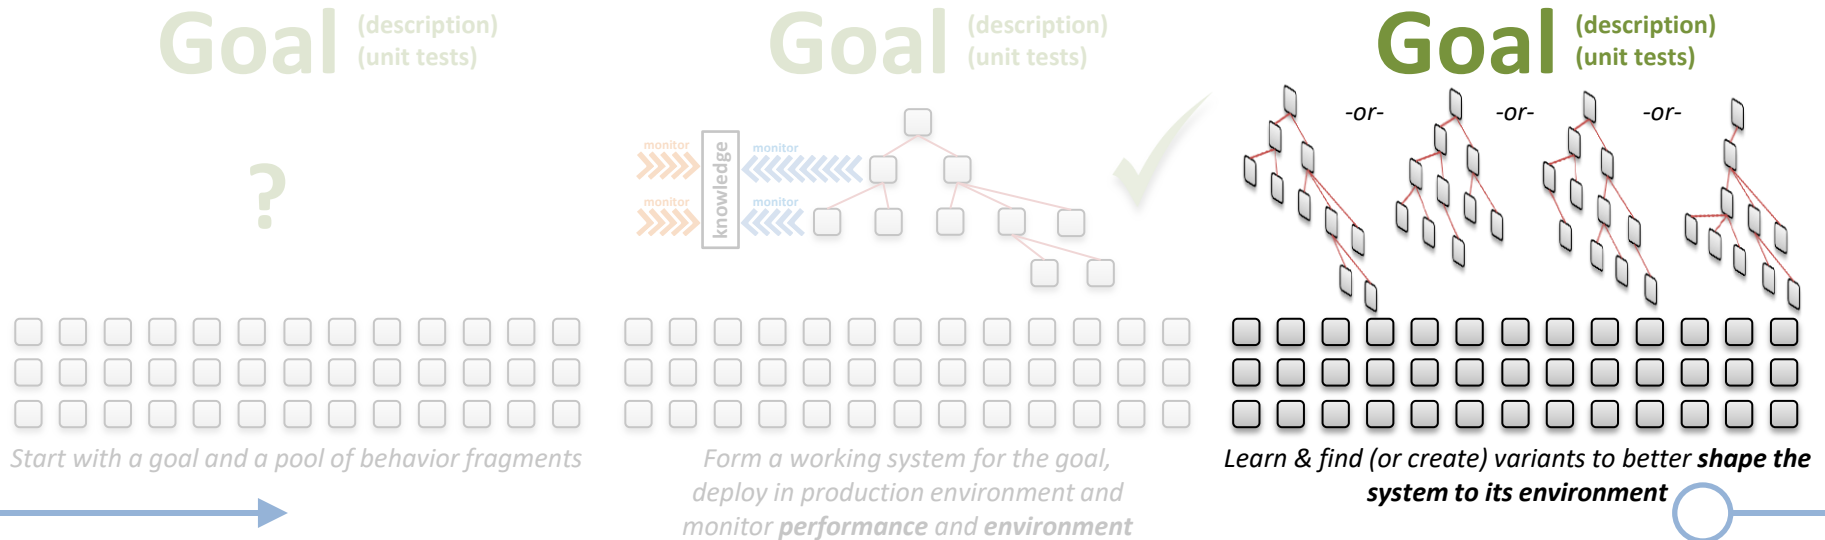


Form a working system for the goal,  
deploy in production environment and  
monitor **performance** and **environment**



# CONCEPT OVERVIEW

- We propose a paradigm of **continuous self-assembly**, in which the *initial construction* of software and its *later adaptation* are one continuous machine-driven process



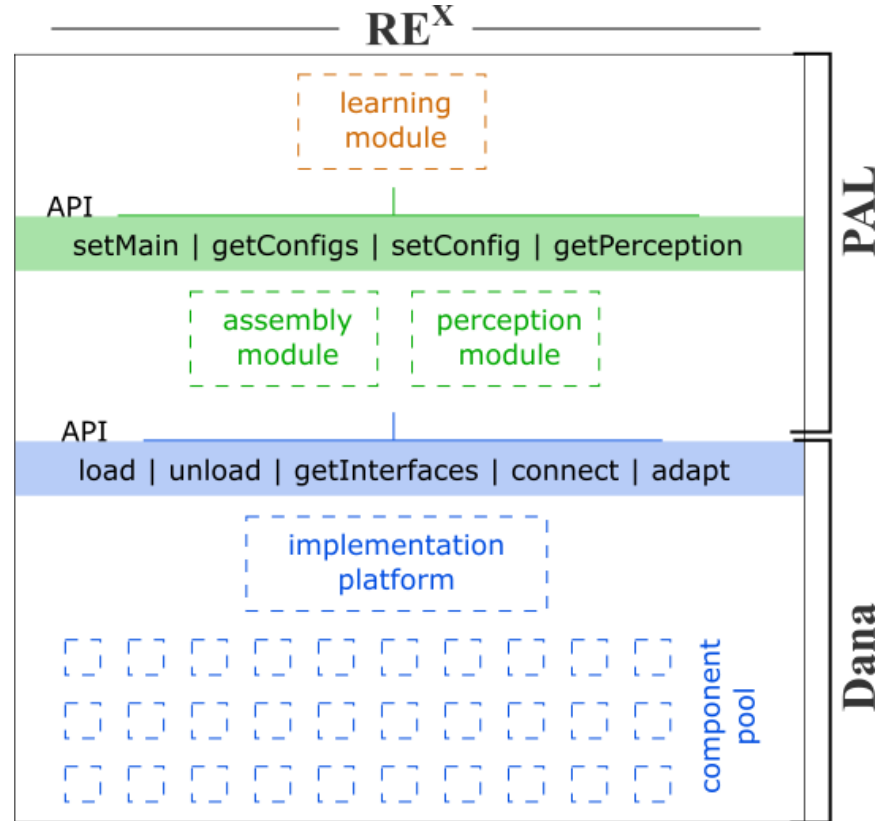
# CONTRIBUTIONS

- **Implementation platform (*Dana*):** A programming language with which to create small software building blocks that can be assembled into emergent systems, with near-zero-cost runtime adaptation for online exploration.
- **Perception, assembly and learning framework (*PAL*):** A framework built with Dana to discover & assemble emergent software, perceive its effectiveness and deployment conditions, and feed perception data to online learning.
- **Online learning approach:** An application of statistical linear bandits, using Thompson sampling, to help solve the search space explosion inherent in our approach, by sharing beliefs about components across possible configurations.

*Example system: an emergent, self-assembling web server*



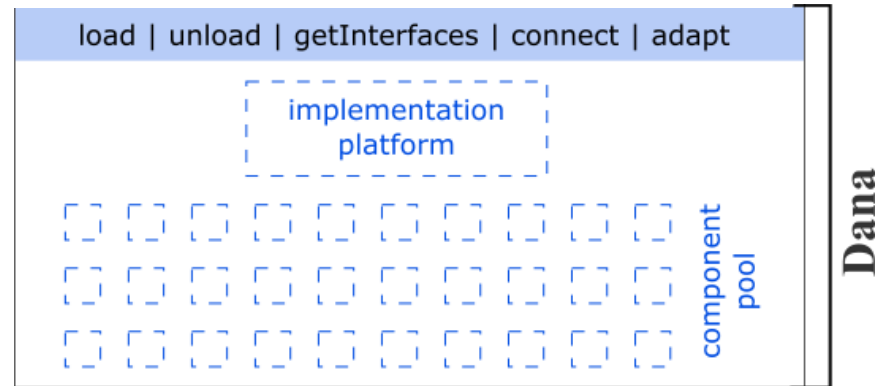
# APPROACH IN DETAIL





# APPROACH IN DETAIL — IMPLEMENTATION PLATFORM

- Uses a **component-based development** paradigm, but:
  - (i) infused in a generalised systems programming language;
  - (ii) supporting very fast, fine-grained runtime adaptation; and
  - (iii) removing the need for wiring diagrams / configurations




# APPROACH IN DETAIL — IMPLEMENTATION PLATFORM

- Uses a **component-based development** paradigm

```
interface RequestHandler {  
    void handleRequest(TCPSocket s)  
}
```

```
component provides App requires net.TCPSocket,  
    net.TCPServerSocket,  
    request.RequestHandler rh {
```

```
    int App:main(AppParam params[]) {  
        TCPServerSocket host = new TCPServerSocket()  
        host.bind(TCPServerSocket.ANY_ADDRESS, 8080)  
  
        while (true) {  
            TCPSocket client = new TCPSocket()  
            if (client.accept(host))  
                rh.handleRequest(client)  
        }  
  
        return 0  
    }  
}
```



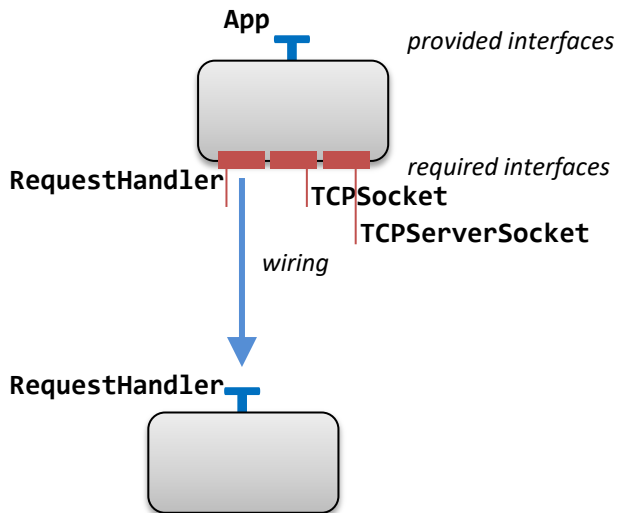
# APPROACH IN DETAIL — IMPLEMENTATION PLATFORM

- Uses a **component-based development** paradigm

```
interface RequestHandler {  
    void handleRequest(TCPSocket s)  
}
```

```
component provides App requires net.TCPSocket,  
    net.TCPServerSocket,  
    request.RequestHandler rh {
```

```
    int App:main(AppParam params[]) {  
        TCPServerSocket host = new TCPServerSocket()  
        host.bind(TCPServerSocket.ANY_ADDRESS, 8080)  
  
        while (true) {  
            TCPSocket client = new TCPSocket()  
            if (client.accept(host))  
                rh.handleRequest(client)  
        }  
  
        return 0  
    }  
}
```



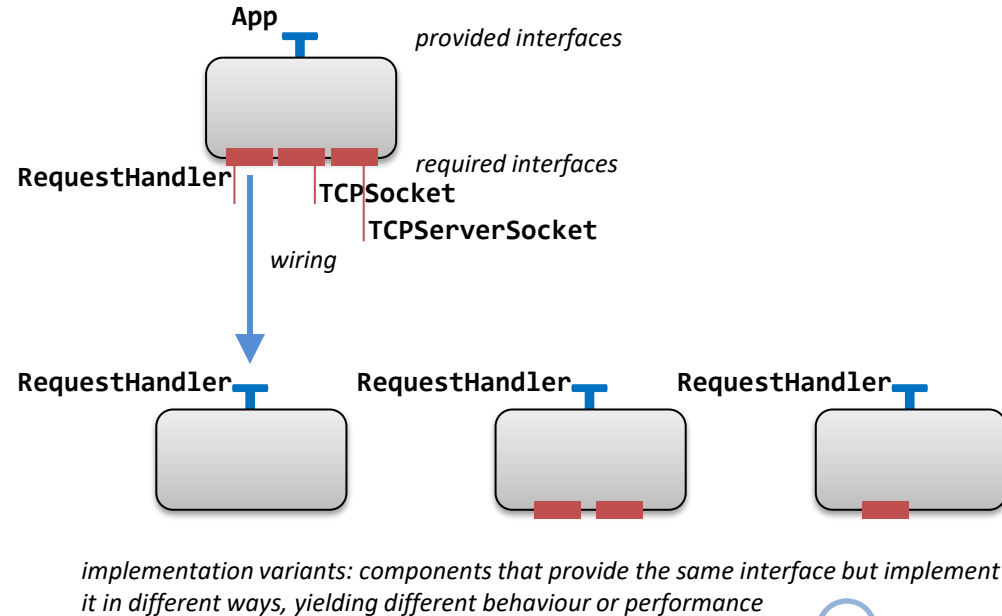
# APPROACH IN DETAIL — IMPLEMENTATION PLATFORM

- Uses a **component-based development** paradigm

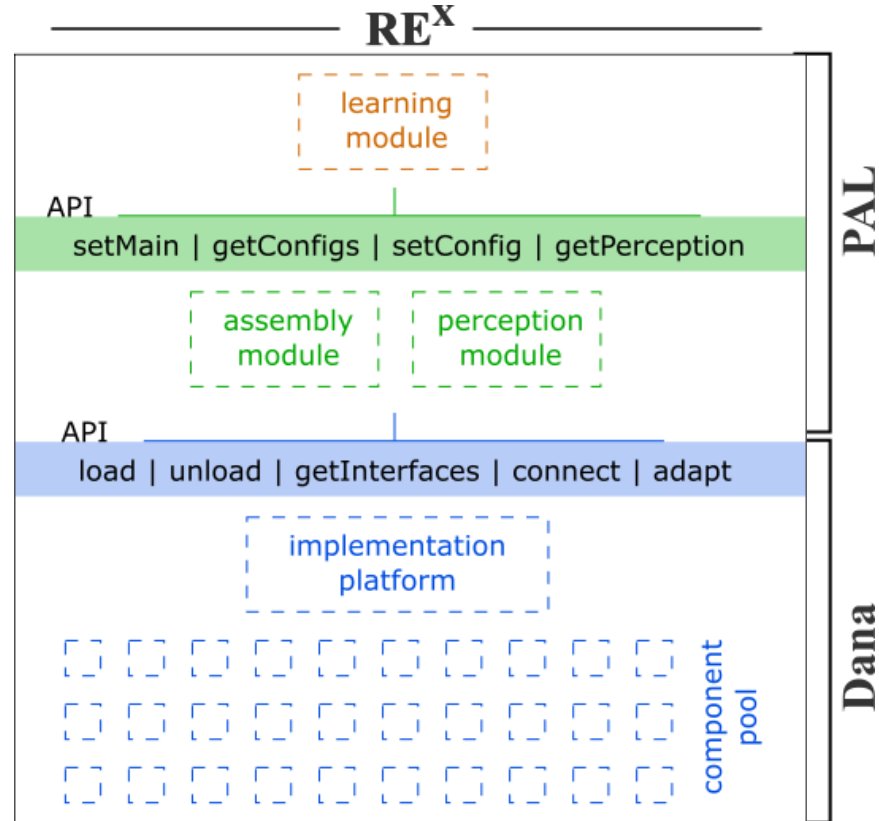
```
interface RequestHandler {  
    void handleRequest(TCPSocket s)  
}
```

```
component provides App requires net.TCPSocket,  
                    net.TCPServerSocket,  
                    request.RequestHandler rh {
```

```
    int App:main(AppParam params[]) {  
        TCPServerSocket host = new TCPServerSocket()  
        host.bind(TCPServerSocket.ANY_ADDRESS, 8080)  
  
        while (true) {  
            TCPSocket client = new TCPSocket()  
            if (client.accept(host))  
                rh.handleRequest(client)  
        }  
  
        return 0  
    }  
}
```

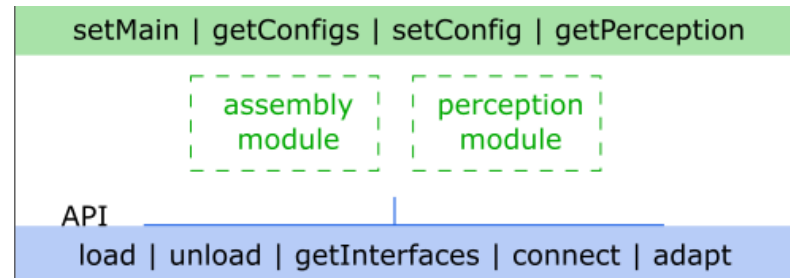


# APPROACH IN DETAIL



# APPROACH IN DETAIL — PERCEPTION, ASSEMBLY AND LEARNING

- A way to **abstract entire software systems** for machine learning into: reward; environment; actions



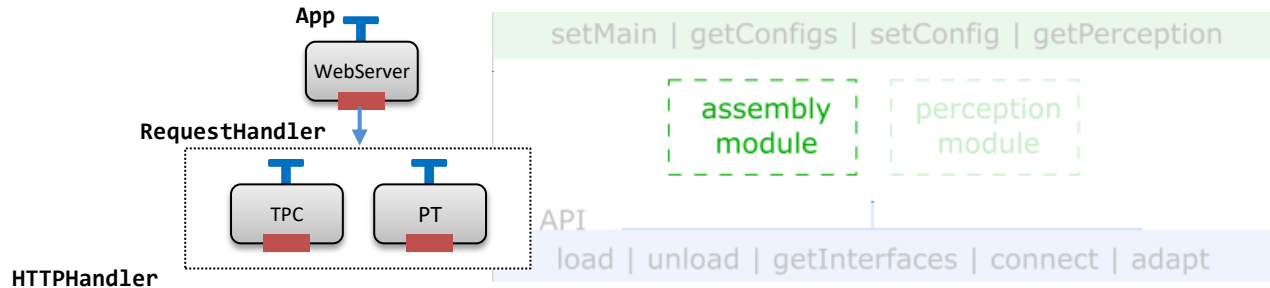
# APPROACH IN DETAIL — PERCEPTION, ASSEMBLY AND LEARNING

- A way to **abstract entire software systems** for machine learning into: reward; environment; actions



# APPROACH IN DETAIL — PERCEPTION, ASSEMBLY AND LEARNING

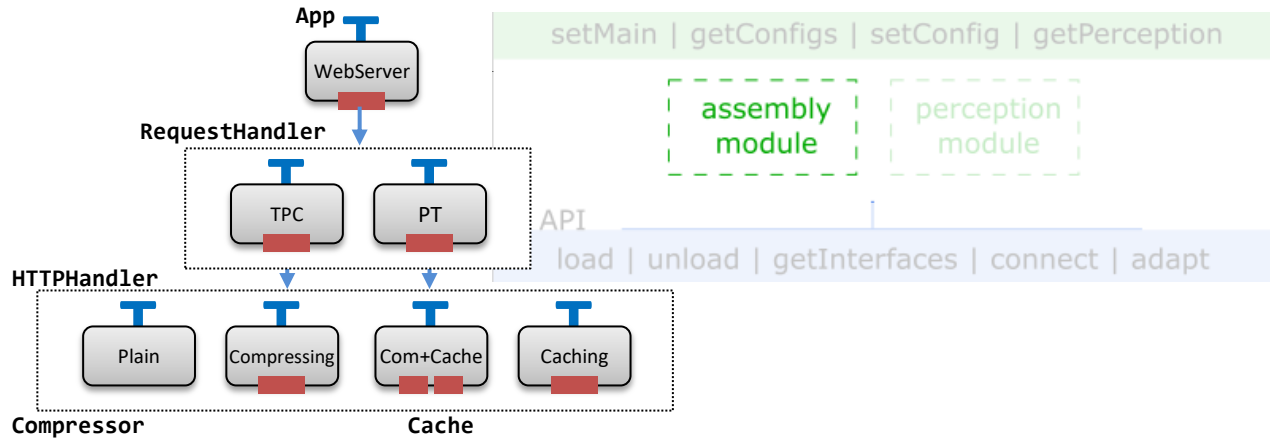
- A way to **abstract entire software systems** for machine learning into: reward; environment; actions





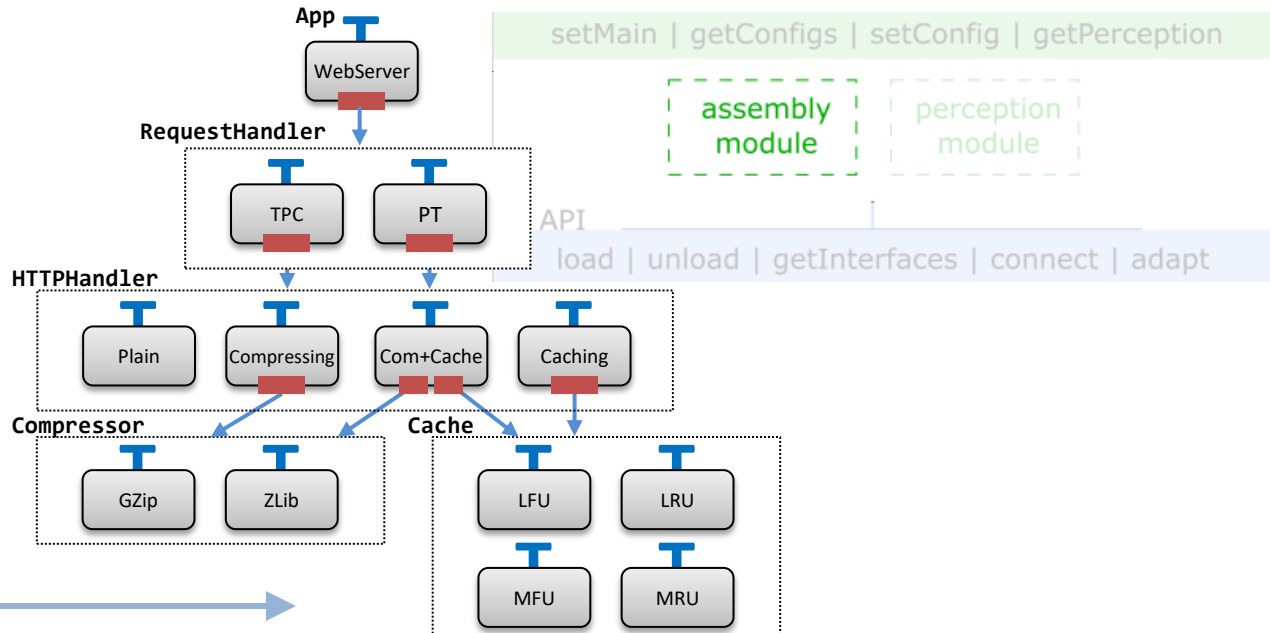
# APPROACH IN DETAIL — PERCEPTION, ASSEMBLY AND LEARNING

- A way to **abstract entire software systems** for machine learning into: reward; environment; actions



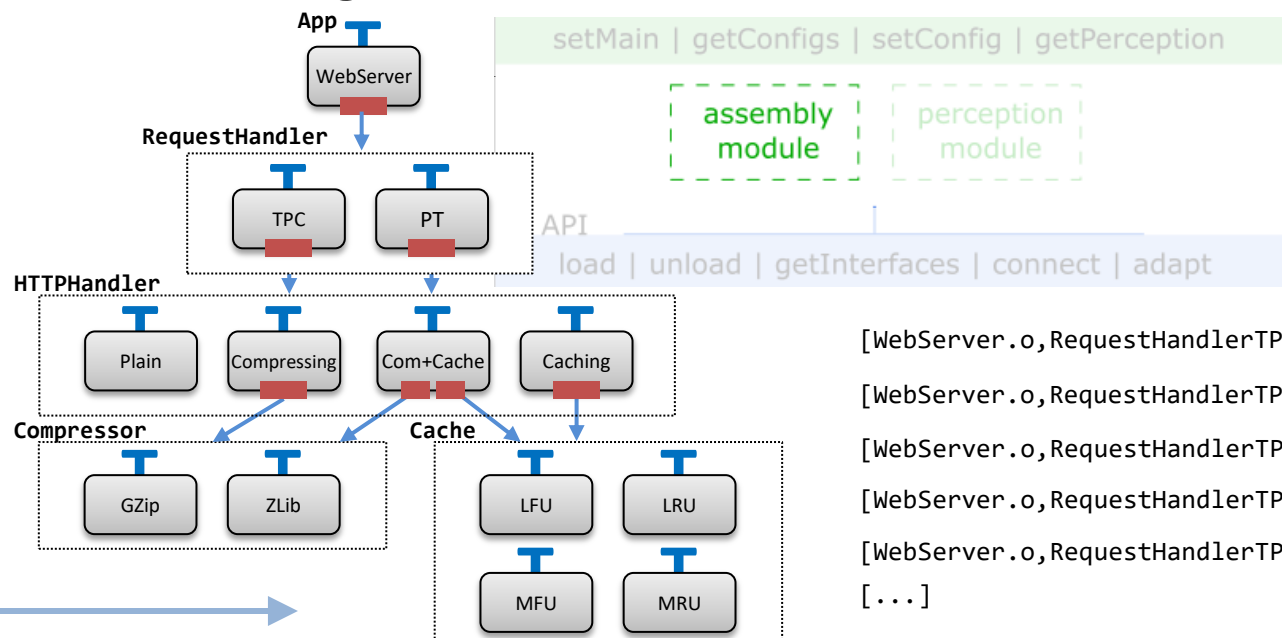
# APPROACH IN DETAIL — PERCEPTION, ASSEMBLY AND LEARNING

- A way to **abstract entire software systems** for machine learning into: reward; environment; actions



# APPROACH IN DETAIL — PERCEPTION, ASSEMBLY AND LEARNING

- A way to **abstract entire software systems** for machine learning into: reward; environment; actions



[WebServer.o,RequestHandlerTPC.o,HTTPHandler.o,TCP.o,File.o,...]

[WebServer.o,RequestHandlerTPC.o,HTTPHandlerZ.o,GZip.o,TCP.o,...]

[WebServer.o,RequestHandlerTPC.o,HTTPHandlerCZ.o,CacheMFU.o,GZip.o,...]

[WebServer.o,RequestHandlerTPC.o,HTTPHandlerCZ.o,CacheLRU.o,GZip.o,...]

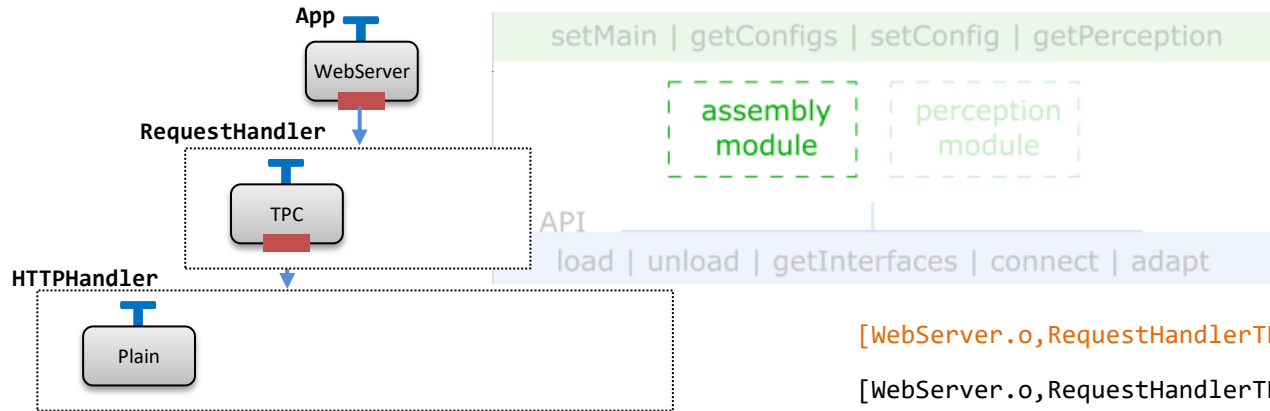
[WebServer.o,RequestHandlerTPC.o,HTTPHandlerC.o,CacheMFU.o,TCP.o,...]

[...]



# APPROACH IN DETAIL — PERCEPTION, ASSEMBLY AND LEARNING

- A way to **abstract entire software systems** for machine learning into: reward; environment; actions



[WebServer.o, RequestHandlerTPC.o, HTTPHandler.o, TCP.o, File.o, ...]

[WebServer.o, RequestHandlerTPC.o, HTTPHandlerZ.o, GZip.o, TCP.o, ...]

[WebServer.o, RequestHandlerTPC.o, HTTPHandlerCZ.o, CacheMFU.o, GZip.o, ...]

[WebServer.o, RequestHandlerTPC.o, HTTPHandlerCZ.o, CacheLRU.o, GZip.o, ...]

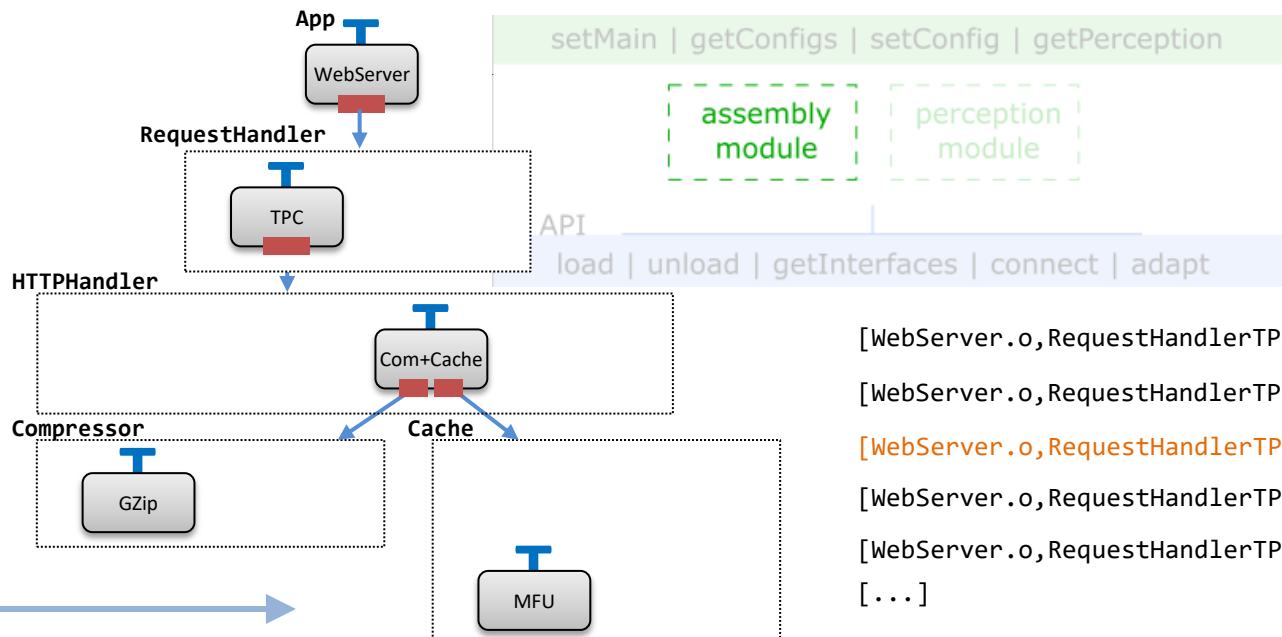
[WebServer.o, RequestHandlerTPC.o, HTTPHandlerC.o, CacheMFU.o, TCP.o, ...]

[...]



# APPROACH IN DETAIL — PERCEPTION, ASSEMBLY AND LEARNING

- A way to **abstract entire software systems** for machine learning into: reward; environment; actions



```
[WebServer.o,RequestHandlerTPC.o,HTTPHandler.o,TCP.o,File.o,...]
```

```
[WebServer.o,RequestHandlerTPC.o,HTTPHandlerZ.o,GZip.o,TCP.o,...]
```

[[WebServer.o](#), [RequestHandlerTPC.o](#), [HTTPHandlerCZ.o](#), [CacheMFU.o](#), [GZip.o](#), ...]

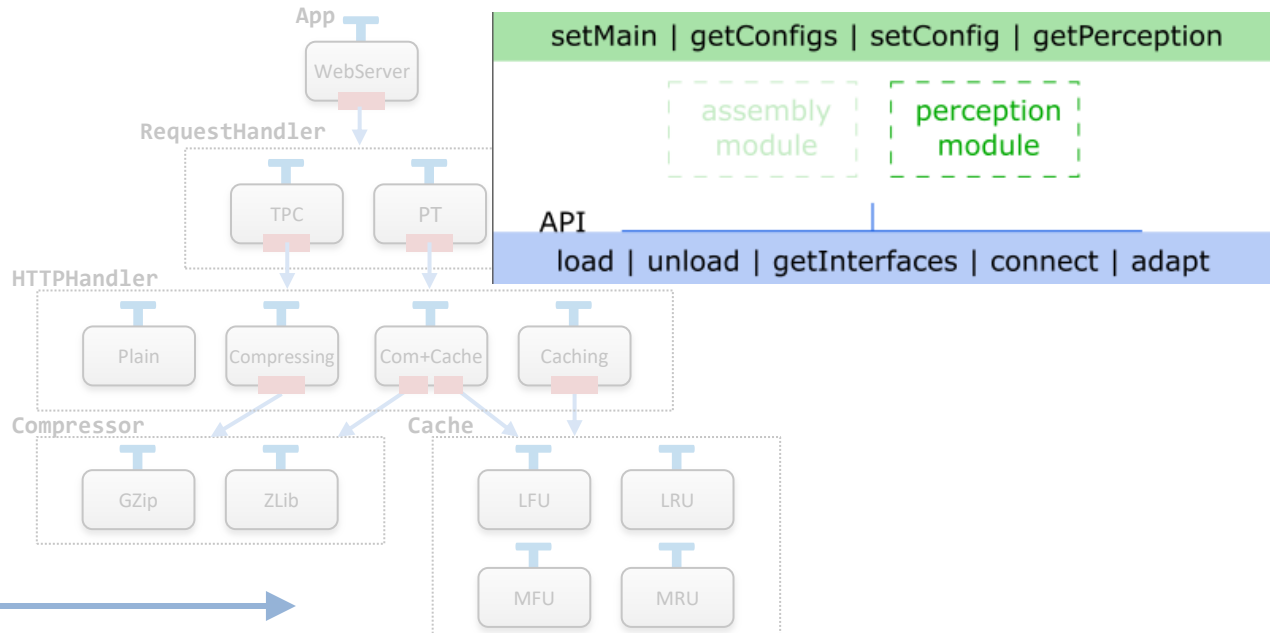
```
[WebServer.o,RequestHandlerTPC.o,HTTPHandlerCZ.o,CacheLRU.o,GZip.o,...
```

```
[WebServer.o, RequestHandlerTPC.o, HTTPHandlerC.o, CacheMFU.o, TCP.o, ...]
```

$$[\dots]$$

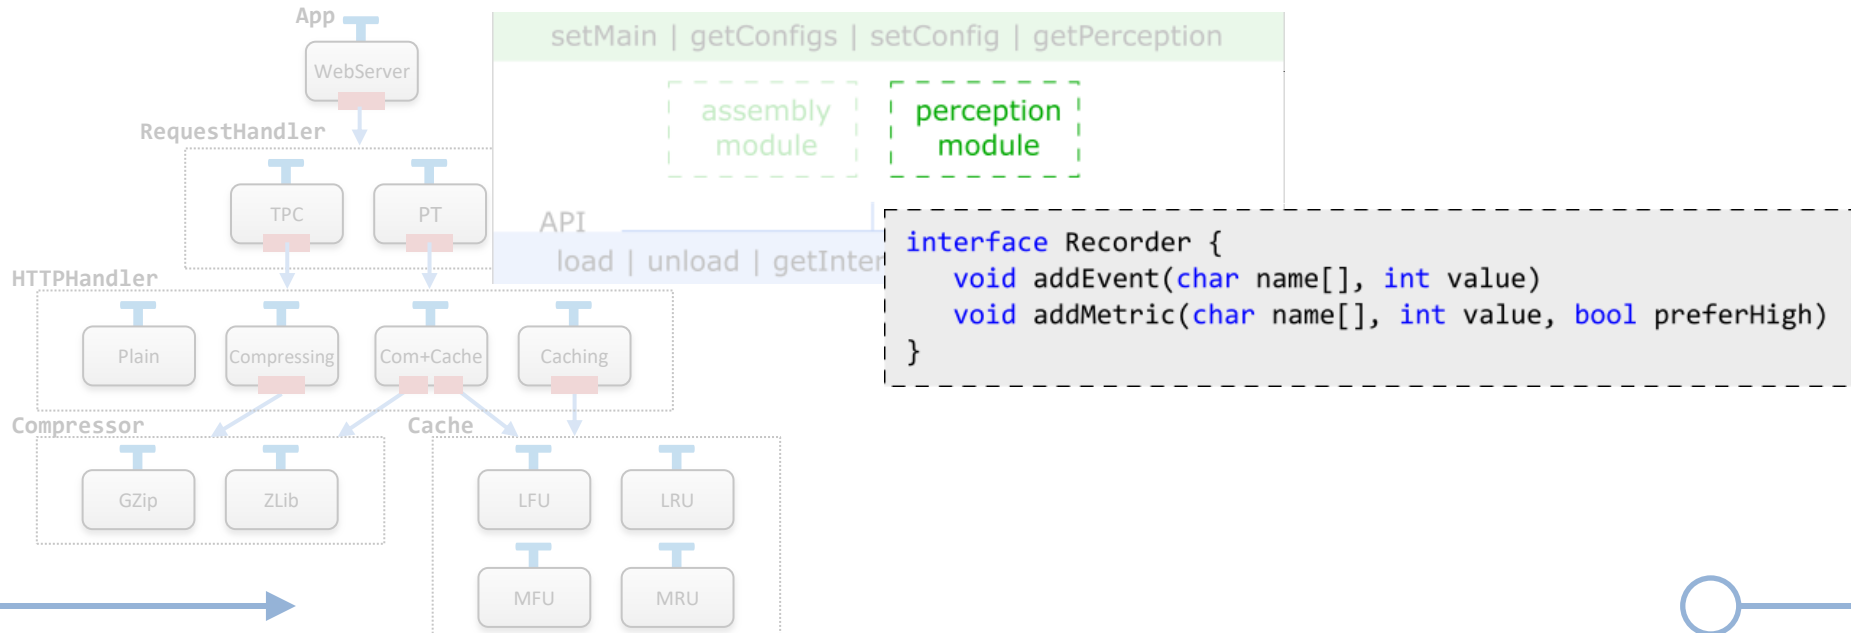
# APPROACH IN DETAIL — PERCEPTION, ASSEMBLY AND LEARNING

- A way to **abstract entire software systems** for machine learning into: reward; environment; actions



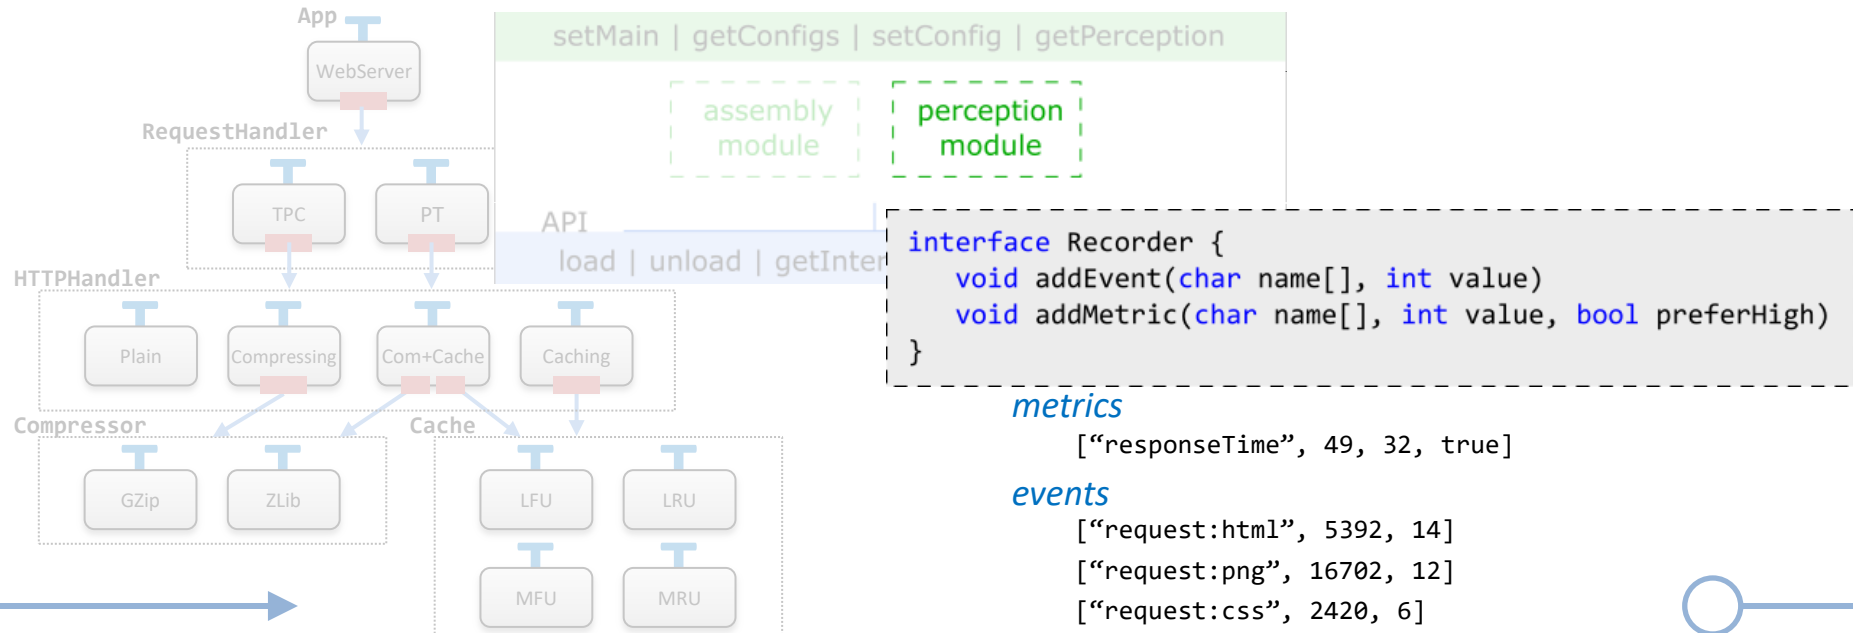
# APPROACH IN DETAIL — PERCEPTION, ASSEMBLY AND LEARNING

- A way to **abstract entire software systems** for machine learning into: reward; environment; actions




# APPROACH IN DETAIL — PERCEPTION, ASSEMBLY AND LEARNING

- A way to **abstract entire software systems** for machine learning into: reward; environment; actions





# APPROACH IN DETAIL



# APPROACH IN DETAIL — LEARNING MODULE

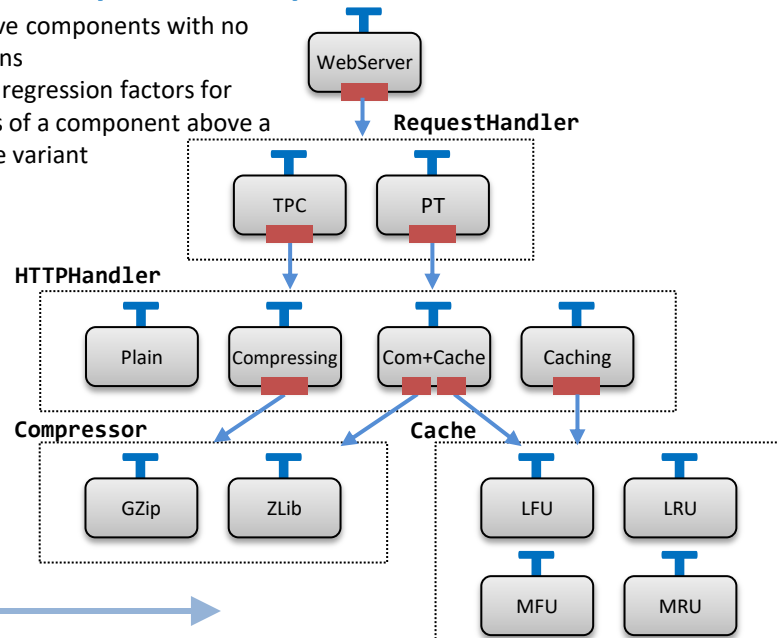
- An **online learning algorithm** which helps solve the *search space explosion*, and balances *exploration* and *exploitation*



# APPROACH IN DETAIL — LEARNING MODULE

- An **online learning algorithm** which helps solve the *search space explosion*, and balances *exploration* and *exploitation*

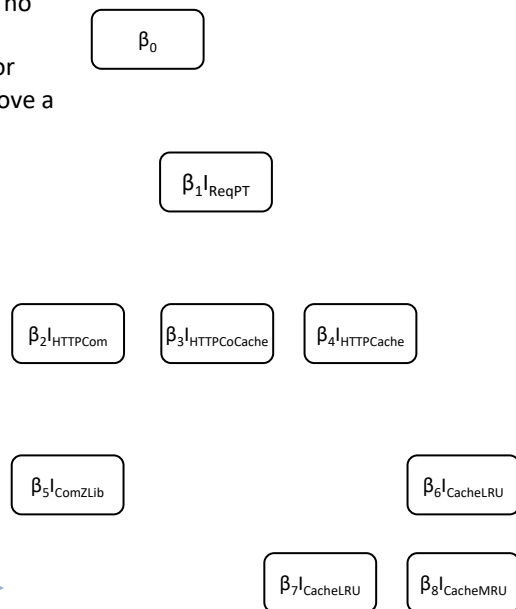
> remove components with no variations  
> make regression factors for variants of a component above a baseline variant



# APPROACH IN DETAIL — LEARNING MODULE

- An **online learning algorithm** which helps solve the *search space explosion*, and balances *exploration* and *exploitation*

> remove components with no variations  
> make regression factors for variants of a component above a baseline variant



# APPROACH IN DETAIL — LEARNING MODULE

- An **online learning algorithm** which helps solve the *search space explosion*, and balances *exploration* and *exploitation*


|                             |             |             |              |              |     |
|-----------------------------|-------------|-------------|--------------|--------------|-----|
| $\beta_0$                   | TPC.Plain   | PT.Plain    | TPC.Com.GZip | TPC.Com.ZLib | ... |
| $\beta_1^{I_{ReqPT}}$       | $\beta_0 1$ | $\beta_0 1$ | $\beta_0 1$  | $\beta_0 1$  | ... |
| $\beta_2^{I_{HTTPCom}}$     | $\beta_1 0$ | $\beta_1 1$ | $\beta_1 0$  | $\beta_1 0$  | ... |
| $\beta_3^{I_{HTTPCoCache}}$ | $\beta_2 0$ | $\beta_2 0$ | $\beta_2 1$  | $\beta_2 1$  | ... |
| $\beta_4^{I_{HTTPCache}}$   | $\beta_3 0$ | $\beta_3 0$ | $\beta_3 0$  | $\beta_3 0$  | ... |
| $\beta_5^{I_{ComZLib}}$     | $\beta_4 0$ | $\beta_4 0$ | $\beta_4 0$  | $\beta_4 0$  | ... |
| $\beta_6^{I_{CacheLRU}}$    | $\beta_5 0$ | $\beta_5 0$ | $\beta_5 0$  | $\beta_5 1$  | ... |
| $\beta_7^{I_{CacheLRU}}$    | $\beta_6 0$ | $\beta_6 0$ | $\beta_6 0$  | $\beta_6 0$  | ... |
| $\beta_8^{I_{CacheMRU}}$    | $\beta_7 0$ | $\beta_7 0$ | $\beta_7 0$  | $\beta_7 0$  | ... |
|                             | $\beta_8 0$ | $\beta_8 0$ | $\beta_8 0$  | $\beta_8 0$  | ... |



# APPROACH IN DETAIL — LEARNING MODULE

- An **online learning algorithm** which helps solve the *search space explosion*, and balances *exploration* and *exploitation*

|                            |             |             |              |              |     |
|----------------------------|-------------|-------------|--------------|--------------|-----|
| $\beta_0$                  | TPC.Plain   | PT.Plain    | TPC.Com.GZip | TPC.Com.ZLib | ... |
| $\beta_1^{I_{ReqPT}}$      | $\beta_0 1$ | $\beta_0 1$ | $\beta_0 1$  | $\beta_0 1$  | ... |
| $\beta_2^{I_{HTTCom}}$     | $\beta_1 0$ | $\beta_1 1$ | $\beta_1 0$  | $\beta_1 0$  | ... |
| $\beta_3^{I_{HTTCoCache}}$ | $\beta_2 0$ | $\beta_2 0$ | $\beta_2 1$  | $\beta_2 1$  | ... |
| $\beta_4^{I_{HTTCCache}}$  | $\beta_3 0$ | $\beta_3 0$ | $\beta_3 0$  | $\beta_3 0$  | ... |
| $\beta_5^{I_{ComZLib}}$    | $\beta_4 0$ | $\beta_4 0$ | $\beta_4 0$  | $\beta_4 0$  | ... |
| $\beta_6^{I_{CacheLRU}}$   | $\beta_5 0$ | $\beta_5 0$ | $\beta_5 0$  | $\beta_5 1$  | ... |
| $\beta_7^{I_{CacheLRU}}$   | $\beta_6 0$ | $\beta_6 0$ | $\beta_6 0$  | $\beta_6 0$  | ... |
| $\beta_8^{I_{CacheMRU}}$   | $\beta_7 0$ | $\beta_7 0$ | $\beta_7 0$  | $\beta_7 0$  | ... |
|                            | $\beta_8 0$ | $\beta_8 0$ | $\beta_8 0$  | $\beta_8 0$  | ... |




every  $n$  seconds, select a composition that we either (1) know little about; or (2) know performs well

# APPROACH IN DETAIL — LEARNING MODULE

- An **online learning algorithm** which helps solve the *search space explosion*, and balances *exploration* and *exploitation*

|                            |             |             |              |              |     |
|----------------------------|-------------|-------------|--------------|--------------|-----|
| $\beta_0$                  | TPC.Plain   | PT.Plain    | TPC.Com.GZip | TPC.Com.ZLib | ... |
| $\beta_1^{I_{ReqPT}}$      | $\beta_0 1$ | $\beta_0 1$ | $\beta_0 1$  | $\beta_0 1$  | ... |
| $\beta_2^{I_{HTTCom}}$     | $\beta_1 0$ | $\beta_1 1$ | $\beta_1 0$  | $\beta_1 0$  | ... |
| $\beta_3^{I_{HTTCoCache}}$ | $\beta_2 0$ | $\beta_2 0$ | $\beta_2 1$  | $\beta_2 1$  | ... |
| $\beta_4^{I_{HTTCCache}}$  | $\beta_3 0$ | $\beta_3 0$ | $\beta_3 0$  | $\beta_3 0$  | ... |
| $\beta_5^{I_{ComZLib}}$    | $\beta_4 0$ | $\beta_4 0$ | $\beta_4 0$  | $\beta_4 0$  | ... |
| $\beta_6^{I_{CacheLRU}}$   | $\beta_5 0$ | $\beta_5 0$ | $\beta_5 0$  | $\beta_5 1$  | ... |
| $\beta_7^{I_{CacheLRU}}$   | $\beta_6 0$ | $\beta_6 0$ | $\beta_6 0$  | $\beta_6 0$  | ... |
| $\beta_8^{I_{CacheMRU}}$   | $\beta_7 0$ | $\beta_7 0$ | $\beta_7 0$  | $\beta_7 0$  | ... |
|                            | $\beta_8 0$ | $\beta_8 0$ | $\beta_8 0$  | $\beta_8 0$  | ... |

select




# APPROACH IN DETAIL — LEARNING MODULE

- An **online learning algorithm** which helps solve the *search space explosion*, and balances *exploration* and *exploitation*

|                             |             |             |              |              |     |
|-----------------------------|-------------|-------------|--------------|--------------|-----|
| $\beta_0$                   | TPC.Plain   | PT.Plain    | TPC.Com.GZip | TPC.Com.ZLib | ... |
| $\beta_1^{I_{ReqPT}}$       | $\beta_0 1$ | $\beta_0 1$ | $\beta_0 1$  | $\beta_0 1$  | ... |
| $\beta_2^{I_{HTTPCom}}$     | $\beta_1 0$ | $\beta_1 1$ | $\beta_1 0$  | $\beta_1 0$  | ... |
| $\beta_3^{I_{HTTPCoCache}}$ | $\beta_2 0$ | $\beta_2 0$ | $\beta_2 1$  | $\beta_2 1$  | ... |
| $\beta_4^{I_{HTTPCache}}$   | $\beta_3 0$ | $\beta_3 0$ | $\beta_3 0$  | $\beta_3 0$  | ... |
| $\beta_5^{I_{ComZLib}}$     | $\beta_4 0$ | $\beta_4 0$ | $\beta_4 0$  | $\beta_4 0$  | ... |
| $\beta_6^{I_{CacheLRU}}$    | $\beta_5 0$ | $\beta_5 0$ | $\beta_5 0$  | $\beta_5 1$  | ... |
| $\beta_7^{I_{CacheLRU}}$    | $\beta_6 0$ | $\beta_6 0$ | $\beta_6 0$  | $\beta_6 0$  | ... |
| $\beta_8^{I_{CacheMRU}}$    | $\beta_7 0$ | $\beta_7 0$ | $\beta_7 0$  | $\beta_7 0$  | ... |
|                             | $\beta_8 0$ | $\beta_8 0$ | $\beta_8 0$  | $\beta_8 0$  | ... |

select



observe reward, update estimates



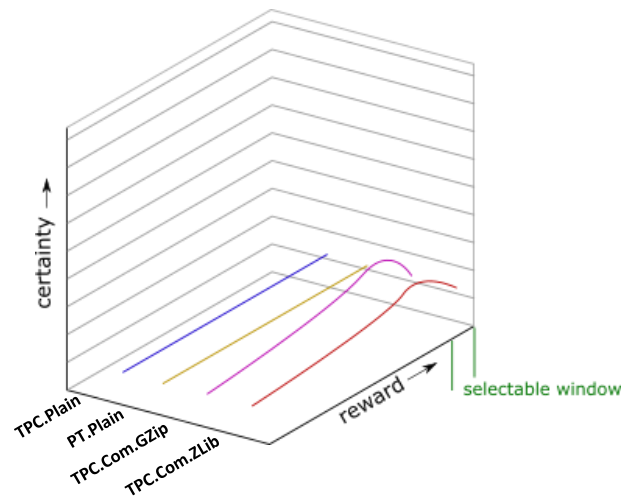


# APPROACH IN DETAIL — LEARNING MODULE

- An **online learning algorithm** which helps solve the *search space explosion*, and balances *exploration* and *exploitation*

|                            |             |             |              |              |     |
|----------------------------|-------------|-------------|--------------|--------------|-----|
| $\beta_0$                  | TPC.Plain   | PT.Plain    | TPC.Com.GZip | TPC.Com.ZLib | ... |
| $\beta_1^{I_{ReqPT}}$      | $\beta_0 1$ | $\beta_0 1$ | $\beta_0 1$  | $\beta_0 1$  | ... |
| $\beta_2^{I_{HTTCom}}$     | $\beta_1 0$ | $\beta_1 1$ | $\beta_1 0$  | $\beta_1 0$  | ... |
| $\beta_3^{I_{HTTCoCache}}$ | $\beta_2 0$ | $\beta_2 0$ | $\beta_2 1$  | $\beta_2 1$  | ... |
| $\beta_4^{I_{HTTCCache}}$  | $\beta_3 0$ | $\beta_3 0$ | $\beta_3 0$  | $\beta_3 0$  | ... |
| $\beta_5^{I_{ComZLib}}$    | $\beta_4 0$ | $\beta_4 0$ | $\beta_4 0$  | $\beta_4 0$  | ... |
| $\beta_6^{I_{CacheLRU}}$   | $\beta_5 0$ | $\beta_5 0$ | $\beta_5 0$  | $\beta_5 1$  | ... |
| $\beta_7^{I_{CacheLRU}}$   | $\beta_6 0$ | $\beta_6 0$ | $\beta_6 0$  | $\beta_6 0$  | ... |
| $\beta_8^{I_{CacheMRU}}$   | $\beta_7 0$ | $\beta_7 0$ | $\beta_7 0$  | $\beta_7 0$  | ... |
|                            | $\beta_8 0$ | $\beta_8 0$ | $\beta_8 0$  | $\beta_8 0$  | ... |

select



share estimates across configurations




# APPROACH IN DETAIL — LEARNING MODULE

- An **online learning algorithm** which helps solve the *search space explosion*, and balances *exploration* and *exploitation*

|                             |             |             |              |              |     |
|-----------------------------|-------------|-------------|--------------|--------------|-----|
| $\beta_0$                   | TPC.Plain   | PT.Plain    | TPC.Com.GZip | TPC.Com.ZLib | ... |
| $\beta_1^{I_{ReqPT}}$       | $\beta_0 1$ | $\beta_0 1$ | $\beta_0 1$  | $\beta_0 1$  | ... |
| $\beta_2^{I_{HTTPCom}}$     | $\beta_1 0$ | $\beta_1 1$ | $\beta_1 0$  | $\beta_1 0$  | ... |
| $\beta_3^{I_{HTTPCoCache}}$ | $\beta_2 0$ | $\beta_2 0$ | $\beta_2 1$  | $\beta_2 1$  | ... |
| $\beta_4^{I_{HTTPCache}}$   | $\beta_3 0$ | $\beta_3 0$ | $\beta_3 0$  | $\beta_3 0$  | ... |
| $\beta_5^{I_{ComZLib}}$     | $\beta_4 0$ | $\beta_4 0$ | $\beta_4 0$  | $\beta_4 0$  | ... |
| $\beta_6^{I_{CacheLRU}}$    | $\beta_5 0$ | $\beta_5 0$ | $\beta_5 0$  | $\beta_5 1$  | ... |
| $\beta_7^{I_{CacheLRU}}$    | $\beta_6 0$ | $\beta_6 0$ | $\beta_6 0$  | $\beta_6 0$  | ... |
| $\beta_8^{I_{CacheMRU}}$    | $\beta_7 0$ | $\beta_7 0$ | $\beta_7 0$  | $\beta_7 0$  | ... |
|                             | $\beta_8 0$ | $\beta_8 0$ | $\beta_8 0$  | $\beta_8 0$  | ... |

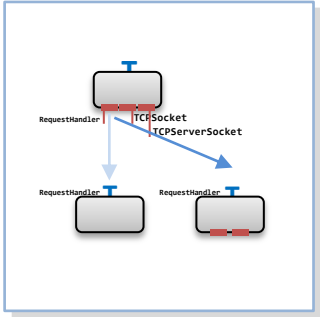
select



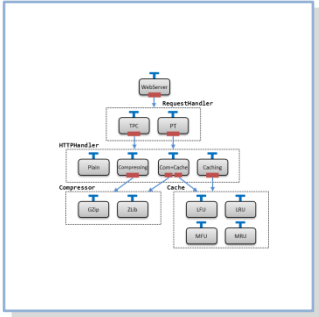
continue learning...



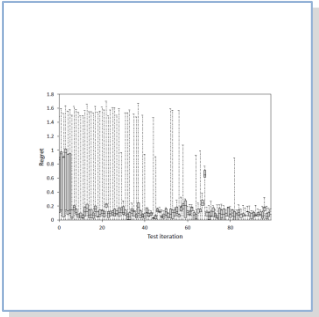
# EVALUATION



Adaptation speed



Performance ground truth



Learning characteristics

Motivation

OSDI 2016

Overview

Contributions

Approach in Detail



[35/41]

REX

Evaluation

b.f.porter@lancaster.ac.uk

Barry Porter



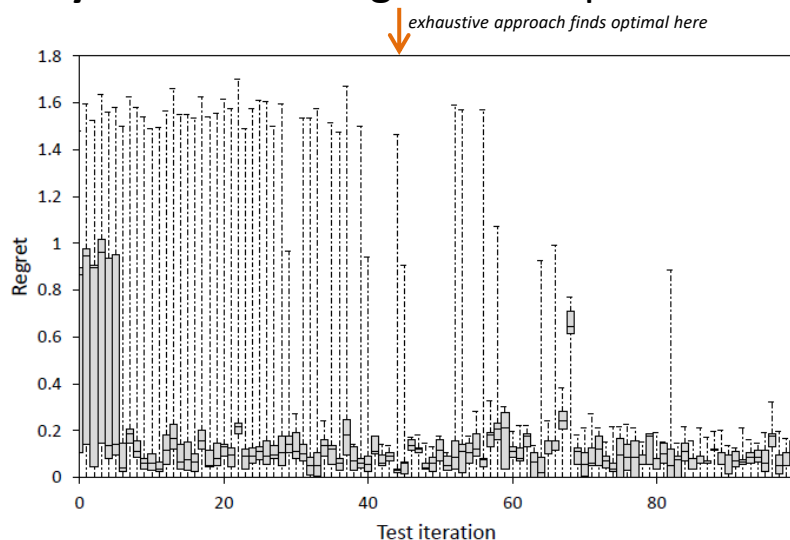
# EVALUATION — LEARNING CHARACTERISTICS

- Using our web server as an example, we evaluate how optimal systems emerge over time in our approach
- We use a set of different workloads (client request patterns), including synthetic and real-world traces
- The only data available for learning is **(i)** the configuration set; and **(ii)** the metrics and events that the system emits (in this case, response times and request types/volumes)




# EVALUATION — LEARNING CHARACTERISTICS

**Key result:** convergence on optimal solution happens much faster than exhaustive search



Workload: small text files




Workload: large image files

Distance from optimal solution over  
time, averaged across 1,000 runs



# EVALUATION — LEARNING CHARACTERISTICS

**Key result:** once learned, workload changes are rapidly detected and adjusted to




Workload: cycling between two different request patterns every 100 seconds

Distance from optimal solution over  
time, averaged across 1,000 runs



# EVALUATION — LEARNING CHARACTERISTICS

**Key result:** convergence occurs in a highly varying real-world workload trace



Workload: real-world workload taken from a publicly available NASA web server trace

Distance from optimal solution over  
time, averaged across 1,000 runs



# EVALUATION — OTHER INSIGHTS

- More broadly, *unexpected* optimal designs that emerged due to machine learning were some of the most interesting results
- This highlighted cases in which our assumptions were wrong about how a given composition would behave, or examples of programmer error / poor design choices






# SUMMARY

- Presented the idea of *emergent software systems* as a new solution to system complexity and deployment dynamics
- We use a paradigm of continuous self-assembly, finding optimal systems via automated composition from small building blocks
- Future work: studying more applications (other server types, AI, robotics); automated generation of variants; automated environment classification; distributed emergent systems (e.g. entire datacentre software landscapes)

→ - download our code at <http://www.projectdana.com> -



Barry  
Porter



Matthew  
Grieves



Roberto  
Rodrigues Filho



David  
Leslie

