



# A SYSTEM TO VERIFY NETWORK BEHAVIOR OF KNOWN CRYPTOGRAPHIC CLIENTS

Andrew Chi, Robert A. Cochran, Marie Nesfield, Michael K. Reiter,  
Cynthia Sturton

University of North Carolina at Chapel Hill

# INVALID COMMAND ATTACKS



## INVALID COMMAND

**Client exhibits behavior, as seen by the server, that is inconsistent with the sanctioned client software.**

### Forms of Exploit:


1. **Maliciously crafted packet**
2. **Valid packets; illegal sequence**

Constrained attacker  
(valid commands only)

**Goal: constrain all attackers**  
to this limited behavior.

# TRANSPORT LAYER SECURITY (TLS) - RFC 5246

- Handshake Protocol
  - Select cipher, authentication, key exchange
- Record Layer
  - Provides confidentiality and integrity
  - Encapsulates other protocols
- Alerts and Heartbeats




**From Jan 2014 to Aug 2016, most of the server-side vulnerabilities in OpenSSL involved invalid commands (23 of 37 required tampering with client behavior).**

# HEARTBLEED (CVE-2014-0160)




- Implementation bug in OpenSSL TLS Heartbeat handler
- Nearly all OpenSSL applications vulnerable for 2 years
- 17% (~500,000) of the Internet's web servers

# HEARTBLEED




# HOW CAN WE DEFEND THE SERVER?



- Behavioral verification: permit authorized client software's behavior only
  - Eliminates entire classes of attack *without knowing about them*
  - Usually requires client modification or sending of client inputs
- Goal: rapid detection of exploit attempts

# BEHAVIORAL VERIFICATION OF A CLIENT




- General case: undecidable
- Specific instances: may be practical
- E.g., detect cheating in online games (Cochran & Reiter 2013)

# SYMBOLIC EXECUTION

```
x = sym_input();  
y = sym_input();  
testme(x, y);
```

```
void testme(int x, int y)  
{  
    int z = 2*y;  
    if (z == x) {  
        if (x > y+10)  
            printf("lol");  
    }  
}
```



Apply SAT solver to obtain concrete test case.


**Example adapted from:** Cristian Cadar, and Koushik Sen. "Symbolic execution for software testing: three decades later." Communications of the ACM 56.2 (2013): 82-90.

# USING SYMBOLIC EXECUTION TO DETECT INVALID COMMAND ATTACKS

```
x = sym_input();  
y = sym_input();  
testme(x,y);  
  
void testme(int x, int y)  
{  
    int z = 2*y;  
    if (z == x) {  
        if (x > y+10)  
            send(z);  
    }  
}
```


Can this program produce...

- $z = 42$ ? Yes ( $x = 42, y = 21$ )
- $z = 41$ ? No ( $z = 2y$  so it must be even)



# CHALLENGES IN VALIDATING CLIENTS IN CRYPTOGRAPHIC PROTOCOLS (1)


- Symbolic execution generally accommodates program variables with unknown values, but their sizes must be known
- Crypto protocols that hide sizes of client-side inputs (e.g., using padding) dramatically grow the search space



- Solution: Explore inputs of different sizes in parallel

# CHALLENGES IN VALIDATING CLIENTS IN CRYPTOGRAPHIC PROTOCOLS (2)

- Some functions are too costly to execute on symbolic inputs
- Example: cryptographic functions
  - AES block cipher is a very complex formula of key and plaintext




- Solution:
  - Give the verifier the session key
  - Defer executing prohibitive functions until inputs can be inferred
    - Any functions not executed then amount to assumptions


# MULTIPASS SYMBOLIC EXECUTION

- Input: user specifies prohibitive functions, using an API
- Algorithm:
  1. Run symbolic execution.
    - a) For each prohibitive function check if any inputs are symbolic
    - b) If so, “skip” the function: return unconstrained symbolic output
    - c) Otherwise, execute the function normally (all inputs are concrete)
  2. Concretize any variables with unique solution
  3. Repeat steps 1-2 until fixed point


# EXAMPLE: TLS CLIENT VALIDATION




# TLS CLIENT VALIDATION PASS 1(A): SYMBOLIC EXECUTION




# TLS CLIENT VALIDATION PASS 1(B): CONCRETIZATION



# TLS CLIENT VALIDATION PASS 2(A): SYMBOLIC EXECUTION



# TLS CLIENT VALIDATION PASS 2(B): CONCRETIZATION



# ASSESSMENT: DETECTING HEARTBLEED (WITHOUT LOOKING FOR IT)


- Malicious s\_client
  - performs handshake
  - sends Heartbleed exploit
- Validation
  - Handshake is verified
  - No explanation found for malicious Heartbeat

Detection in ~2s



```
19:33:58 | CV: Opened socket log "/playpen/bu
19:33:58 | CV: BasicBlock count: 61686
19:33:58 | CV: Creating stage from add_state(
)* @_user_main to i32 (i32, i8**, i8**)*, i
19:33:58 | KLEE: Attempting to open: /home/ac
19:33:59 | KLEE: Attempting to open: /playpen
19:33:59 | KLEE: Attempting to open: /home/ac
19:33:59 | KLEE: Attempting to open: /home/ac
19:33:59 | KLEE: Attempting to open: /playpen
19:33:59 | KLEE: Attempting to open: /home/ac
19:34:00 | CV: Thread 1 executed 7833620 ins
19:34:00 | CV: Generating SearcherStage graph
19:34:00 | CV: Verifier Result: failure (1)
total instructions 7833620
```


# MEASURING PERFORMANCE



# PERFORMANCE EVALUATION

- 21 TLS 1.2 sessions from 3 min. of Gmail activity
- OpenSSL & BoringSSL command line clients
- Single-core verifier (3.2 GHz)
- Cost: 49ms per TLS record
- Lag: median 0.85s, max 15s

NOTE: without server-to-client appdata packets



# OTHER EVALUATION MEASURES

- Parallelization / Stress Test
  - TLS 1.2 + up to 128 bytes of padding (from draft TLS 1.3)
  - 16-thread verifier keeps pace
- Invalid command attack: valid packets, illegal sequence
  - CVE-2015-0205 client authentication vulnerability
  - Verifier rejects attack traffic
- Confirm appropriateness of command line client
  - Unmodified Chrome browser interacting with Apache server
  - Verified using BoringSSL command line client

# SUMMARY

- Behavioral verification for cryptographic clients
  - Multipass symbolic execution handles cryptographic functions
  - Parallelization optimizes search of large state spaces
- Detection of previously unknown client misbehavior
  - E.g., a Heartbleed exploit with no Heartbleed-specific configuration
- Performance roughly keeps pace with real workload
  - Behavioral verification on Gmail TLS sessions