

Brief Announcement: Object Oriented Consensus

Yehuda Afek

School of Computer Science, Tel-Aviv University
afek@tau.ac.il

Edo Cohen

School of Computer Science, Tel-Aviv University
edocohen@tau.ac.il

James Aspnes

School of Computer Science, Yale University
james.aspnes@gmail.com

Danny Vainstein

School of Computer Science, Tel-Aviv University
danvainstein@tau.ac.il

ABSTRACT

We suggest a template that reveals the structure of many consensus algorithms as a generic procedure. The template builds on a new object, *vacillate-adopt-commit* which is an extension of the well known *adopt-commit* object. In addition we extend Aspnes's *conciliator* object to a new object that we call a *reconciliator*. The consensus algorithm template works in rounds of alternating *vacillate-adopt-commit* and *reconciliator* operations. The *vacillate-adopt-commit* object observes the processors' preferences and suggests a preference output with a measure of confidence (*vacillate, adopt or commit*) on the preference. The *reconciliator* ensures termination, by providing new preferences for the processors. We show how several key consensus algorithms exactly fit our template. Here we demonstrate the decomposition of Ben-Or's randomized algorithm. The decomposition of the Phase King Byzantine and the Paxos algorithm are given in the full paper [1]. We analyze and compare our template based on *vacillate-adopt-commit* and *reconciliator* objects to previous work [3, 5], suggesting a decomposition of consensus based on *adopt-commit* and *conciliator* objects. We claim that the three return values of *vacillate-adopt-commit* more accurately describe existing algorithms.

KEYWORDS

Distributed Algorithms; Consensus; Adopt-Commit;

1 INTRODUCTION

The consensus problem, introduced by Lamport, Pease and Shostak [6] resides at the heart of many distributed algorithms such as leader election, database transaction handling, resource allocation, ensuring storage replicas are mutually consistent and many more.

In the consensus protocol between n processors, each with an input value, processors agree on a single common output which was the input to one of them. While consensus is trivial in a non-faulty synchronous environment, it is often more difficult in practice as most distributed networks are asynchronous and must be resilient to faults of various types.

Gafni [5] proposed the *adopt-commit* object, an object fulfilling weaker guarantees than consensus, as a building block of consensus. Aspnes [3] further provided a detailed decomposition of consensus into *adopt-commit* and a complementary *conciliator* object which together form a generic framework describing consensus. In this work we extend these decompositions, into *vacillate-adopt-commit* and *reconciliator* objects, which describe the core of many existing consensus algorithms more accurately than the previous decomposition.

The paper is structured as follows. In Section 3, the consensus template is presented. In Section 4, we demonstrate how the decomposition applies to Ben-Or's consensus algorithm (in our full paper [1] we further apply our decomposition to the Paxos and Phase-King consensus algorithms). Section 5 explores the relation between *vacillate-adopt-commit* and *adopt-commit*, showing the latter is slightly weaker.

2 PRELIMINARIES

Consensus guarantees the following three properties: **Validity** - The output value, u , was provided as an input to the object by some processor. **Termination** - The output is returned within a finite number of steps. **Agreement** - All processors invoking the object receive the same output value u .

Adopt-commit is a weaker form of consensus which guarantees *validity* and *termination*. The object supports a single operation, *adopt-commit*(v) and returns some preference u with a level of confidence, *adopt* or *commit*. The object does not guarantee *agreement*, instead, *adopt-commit* holds the following guarantees: **Coherence** - If some processor receives (*commit*, v), any other returned value (with either *commit* or *adopt*) must have the same value v . **Convergence** - If all processors invoke *adopt-commit*(v) with the same value v , the object must return (*commit*, v) with the same value v to all processors.

Vacillate-adopt-commit is an extension of *adopt-commit* which supports a single operation, *vacillate-adopt-commit*(v) and returns some preference u and a level of confidence, *vacillate*, *adopt* or *commit*. As opposed to *adopt-commit*, *vacillate-adopt-commit* adds a level of confidence and guarantees *validity*, *termination* and *convergence*. **Coherence** is modified to suit three levels of confidence, therefore, *vacillate-adopt-commit* guarantees **coherence over adopt & commit** - If any processor receives (*commit*, v), then any other returned value is either (*commit*, v) or (*adopt*, v) (with the same value v). **Coherence over vacillate & adopt** - If no processor received *commit* and some processor received (*adopt*, u), then every other processor receives either (*adopt*, u) or (*vacillate*, w) where w may be any value.

Permission to make digital or hard copies of part or all of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for third-party components of this work must be honored. For all other uses, contact the owner/author(s).
PODC'17, , July 25-27, 2017, Washington, DC, USA.
© 2017 Copyright held by the owner/author(s).
ACM ISBN 978-1-4503-4992-5/17/07.
<http://dx.doi.org/10.1145/3087801.3087867>

The **reconciliator** object also supports a single operation, *reconciliator*(v) and only returns a preference u . The object guarantees that every value has a non-zero probability (that may depend on n , the number of processors) to be returned to all processors. In the *Byzantine* setting, the *reconciliator* guarantees to return the same value to all processors eventually (i.e., if called sufficiently many times).¹

3 THE GENERIC FORM OF CONSENSUS

The thesis of this paper is that many well known consensus algorithms have the same basic structure consisting of two objects, *VAC*² that checks whether consensus has been reached or not and *reconciliator* that shakes up the preferences of the processors in case of a stalemate. We do not claim that this structure is necessarily better than using *AC* and *conciliators*, but that it more accurately reflects the existing structure of algorithms in the literature.

Informally, the generic consensus algorithms work in rounds. In each round, first the *VAC* is invoked to observe the system state and tell whether consensus has been reached. *VAC* returns to each processor one of three possible outputs: (1) (**commit**, v) which indicates that the system has reached an agreement on value v , (2) (**adopt**, v) which indicates that it is possible that some processors in the system have agreed on the value v , and (3) (**vacillate**, v) indicating that the system is in an indecisive state.

If a processor receives (**commit**, v), it is guaranteed that no other processor receives a *vacillate* value and all outputs return with the same value, v . A processor receiving (**adopt**, v) is guaranteed that any other processor either received a *vacillate* value or received the same preference. Finally, if a processor receives (**vacillate**, v), the only guarantee it has is that no other processor received a *commit* value.

We note the key difference between the *VAC* and *AC* objects. An *adopt-commit* object always returns a new value to be adopted by a process, but this is not consistent with the structure of many consensus protocols in the literature. Adding a third option, i.e., *vacillate*, accounts for situations where the algorithm does not force a process to update its preference.

The question is how termination of the consensus can be guaranteed if the collection of preferences is balanced and the *VAC* continually returns *vacillate*. For that purpose, the *reconciliator* is used to give each *vacillating* processor a new preference with a guarantee to provide a deciding set of preferences with some probability. That is, whenever a processor receives a *vacillate* value from the *VAC* object, the distributed *reconciliator* provides each processor with an alternate preference such that eventually enough processors will get the same preference leading to *VAC* eventually observing agreement. Pseudocode for the consensus template is given in Algorithm 1.

Note that *INIT* is a void function unless stated otherwise. Furthermore, note that the operation, *decide* σ , is followed by a halt

```

1 Consensus( $v$ )
2    $m \leftarrow 0$ 
3   INIT()
4   while true do
5      $m \leftarrow m + 1$ 
6      $(X, \sigma) \leftarrow VAC(v, m)$ 
7      $\sigma' \leftarrow reconciliator(\sigma, m)$ 
8     switch  $X$  do
9       case vacillate:  $v \leftarrow \sigma'$ 
10      case adopt:  $v \leftarrow \sigma$ 
11      case commit:  $v \leftarrow \sigma$  and decide  $\sigma$ 
12    endsw
13  end

```

Algorithm 1: Consensus Template

operation, that is, the processor will decide upon its value and return. The argument m is the phase of the consensus process.

Next we prove that the template indeed achieves consensus, using the *VAC* and *reconciliator* properties.

LEMMA 3.1. *Algorithm 1 is a correct consensus algorithm.*

PROOF. **Agreement** and **validity** follow from *VAC*'s *coherence* and *validity* respectively. **Termination** follows from the *convergence* property of the *reconciliator* object. \square

4 DECOMPOSING BEN-OR'S ALGORITHM

In this section we show how Ben-Or's algorithm [4] can be described using our consensus template. Throughout this section the settings are asynchronous, message-passing model and the number of tolerated crash failures, t , is strictly smaller than $n/2$. Algorithms 2 and 3 are the *VAC*'s and *reconciliator*'s implementations, Lemmas 4.1 and 4.2 prove the implementations correctness, i.e., that they uphold the objects' guarantees.

```

1 Reconciliator( $X, \sigma, m$ )
2   return CoinFlip()

```

Algorithm 2: Ben-Or's *reconciliator* implementation

LEMMA 4.1. *Algorithm 2 is a correct reconciliator implementation.*

PROOF. **Probabilistic convergence**: For any given value v , \mathbb{P} (all processors finish with v after the coin flip) $\geq \frac{1}{2^n} > 0$. \square

LEMMA 4.2. *Algorithm 3 is a correct vacillate-adopt-commit implementation.*

PROOF. The proof is similar to the Ben-Or algorithm correctness proof found in the survey of Aspnes [2].

Validity, termination and **convergence** follow easily from the implementation.

Coherence over commit and adopt, assume a process received (*commit*, v). Therefore, it received more than t *ratify* messages, meaning a non-faulty processor sent out a *ratify* message to all

¹We note that the *reconciliator* weakens Aspnes' *conciliator* [3] definition in that it does not require *validity* (i.e., that the *reconciliator*'s returned value equals some processor's input). This is due to the fact that if all (non-Byzantine) processors' inputs are equal, the *VAC* is defined such that no processor will, anyhow, reach the *reconciliator*.

²Throughout the paper we abbreviate *vacillate-adopt-commit* as *VAC* and *adopt-commit* as *AC*.

```

1 VAC( $v, m$ )
2   send  $\langle 1, v \rangle$  to all
3   wait to receive  $n - t \langle 1, * \rangle$  messages
4   if received more than  $n/2 \langle 1, v \rangle$  messages then
5     | send  $\langle 2, v, \text{ratify} \rangle$  to all
6   else
7     | send  $\langle 2, ? \rangle$  to all
8   end
9   wait to receive  $n - t \langle 2, * \rangle$  messages
10  if received more than  $t \langle 2, v, \text{ratify} \rangle$  messages then
11    | return ( $\text{commit}, v$ )
12  else if received a  $\langle 2, v, \text{ratify} \rangle$  message then
13    | return ( $\text{adopt}, v$ )
14  else
15    | return ( $\text{vacillate}, v$ )
16  end

```

Algorithm 3: Ben-Or's *vacillate-adopt-commit* implementation

other processors - therefore it is enough to show that every 2 *ratify* messages have the same value, v (since then all processes received atleast one *ratify* message with the same value, v). Assume towards contradiction that this isn't the case - meaning messages $\langle 2, v, \text{ratify} \rangle$ and $\langle 2, u, \text{ratify} \rangle$ where $u \neq v$ had been sent. By the first *if* statement, this means there's a processor that sent out both u and v which is a contradiction.

Coherence over adopt and vacillate, since every 2 *ratify* messages have the same value, v , coherence follows. \square

5 VACILLATE-ADOPT-COMMIT VS ADOPT-COMMIT

5.1 Adopt-Commit is Not Enough

The concept of decomposing consensus into separate objects is by no means original and was formally presented in [5]. Later work by Aspnes [3] described a framework of *adopt-commit* objects that detect agreement, and *conciliators* that ensure agreement with some probability. We argue that this decomposition breaches abstraction responsibilities. The consensus algorithm may be viewed as an iterative process which begins in some arbitrary state and terminates with an agreement upon a *valid* value. Under this interpretation, the framework of repetitive *adopt-commit* followed by *conciliator* fails to distinguish between the phases of the consensus procedure in which part of the participants have terminated and others have not. Due to this limitation and the *coherence* constraint, the *conciliator* must be aware of the global state of the consensus procedure in order to ensure *validity*.

In order to make our argument more concrete, we demonstrate how Ben-Or's consensus algorithm cannot be described by a sequence of *adopt-commit* alternating with *conciliator*, while it is naturally described as a sequence of repetitive *vacillate-adopt-commit* followed by *reconciliator*.

To demonstrate the problem with formulating Ben-Or's consensus protocol using a series of *adopt-commit* and *conciliator* objects,

$$U = A_{-1}; A_0; C_1; A_1; C_2; A_2; \dots$$

(where A, C denote *adopt-commit* and *conciliator* objects respectively), consider each round of Ben-Or's algorithm [4]³. Let P be a processor participating in the agreement process. P experiences one of three possible outcomes: (1) not receiving any *ratify* message. (2) receiving up to t *ratify* messages. (3) receiving more than t *ratify* messages.

These outcomes correspond to vacillate, adopt, and commit, respectively. Option 1 fits a processor which received vacillate as it has no guarantees about other values received by other processors. Option 2 corresponds to *adopt* under the *VAC* framework, since by *coherence*, any processor that received (*adopt*, v) is guaranteed that every other processor that received either *vacillate* or *commit*, also received the value v . Option 3 corresponds to *commit*, since any processor that received (*commit*, v) is guaranteed that all other processors received either (*commit*, v) or (*adopt*, v). However, using only *adopt-commit* objects is not enough in order to describe these three options.

It might be tempting to assume that two consecutive *adopt-commit* objects might resolve this entanglement as we have shown that *VAC* may be implemented using two *AC* objects.⁴ We argue this is not the case, that is, we claim that the sequence of $U = A_{-1}; A_0^0; A_0^1; C_1; A_1^0; A_1^1; C_2; \dots$ also fails to describe Ben-Or's consensus protocol. In order to describe option (2) the first *adopt-commit* must return *adopt* while the second returns *commit*. However, the decomposition framework described in [3] requires that upon reception of *commit* the processor immediately decides on the value received, whereas it is possible that in Ben-Or's protocol such a state is reached with value u but a final agreement is achieved with value $u' \neq u$.

5.2 Relation Between Vacillate-Adopt-Commit and Adopt-Commit

In addition to the former section, we demonstrate in the full version of our paper [1] that one may easily simulate an *adopt-commit* object using a *VAC* object. On the other hand, one may successfully simulate *VAC* using two sequential calls to *adopt-commit* objects. In our opinion this further shows an inherent difference between the two objects.

REFERENCES

- [1] Yehuda Afek, James Aspnes, Edo Cohen, and Danny Vainstein. Object oriented consensus. In preparation, February 2017.
- [2] James Aspnes. Randomized protocols for asynchronous consensus. *Distributed Computing*, 16(2-3):165–175, September 2003.
- [3] James Aspnes. A modular approach to shared-memory consensus, with applications to the probabilistic-write model. *Distributed Computing*, 25(2):179–188, May 2012.
- [4] Michael Ben-Or. Another advantage of free choice: Completely asynchronous agreement protocols (extended abstract). In *Proceedings of the Second Annual ACM SIGACT-SIGOPS Symposium on Principles of Distributed Computing, Montreal, Quebec, Canada, August 17-19, 1983*, pages 27–30, 1983.
- [5] Eli Gafni. Round-by-round fault detectors (extended abstract): unifying synchrony and asynchrony. In *Proceedings of the seventeenth annual ACM symposium on Principles of distributed computing*, pages 143–152. ACM, 1998.
- [6] Marshall Pease, Robert Shostak, and Leslie Lamport. Reaching agreement in the presence of faults. *Journal of the ACM (JACM)*, 27(2):228–234, 1980.

³We note that our description follows the presentation of Ben-Or's algorithm given in the survey paper of Aspnes [2]

⁴See [1] for detailed arguments.